

Safety requirements	Visitation allocation	Constrained projection	Budget prices	Policy learning
---------------------	-----------------------	------------------------	---------------	-----------------

Soft budgets

Costly visits

$b = B / L$

Total cost

Hard events

$\times$

Zero capacity

Delete edges

Not priced

Visitation allocation

Width = mass

Reward directs useful visits

Flow balance links decisions

Constrained projection

Reward + KL

Flow balance

Soft total-cost budgets

Hard support mask

Target flow q^*

Budget prices

Graph λ^*

Clip

Filter

+

Applied λ

Feedback β

Actual costs

$J - B$

Policy learning

PPO / TRPO

Prices enter the advantage; flow is a target.

Priced advantage


Actor update

New rollouts

Priced advantage

$$\tilde{A} = (A_r - \sum_k \lambda_k A_{c,k}) / (1 + \sum_k \lambda_k)$$

→ actor



Conditional price limits $\epsilon \rightarrow 0$: LP dual price. Large ϵ : information-projection price.

Clip $\rightarrow$ saturated active-set price; realized-cost switch only with matching