




(a) Learning-time violation



(b) Final return



(c) All-task learning curves



Solved prices reduce learning-time violation versus Lagrangian baselines across both suites. (a,b) Right favors SafeOT-Dual: task-median V ratios (reversed log axis) and normalized final-return differences, with 95% bootstrap intervals. Circles denote SafePO; diamonds, Bullet. Counts show tasks with lower V. Dashes mean not run in that suite; n/r means not reported. † Final task-mean cost exceeds B on 4/6 TRPO-Lag and 5/6 PPO-Lag tasks. (c) Lines are task medians of seed means; blue/purple bands show task IQR. Budget equals 1; Bullet cost uses symlog. Normalizers, seed coverage and per-task curves are in the appendix.