


# SafeOT execution scale: measured reward / risk trade-offs


Each summary uses training seeds 1, 2, 3 and the same 3 x 10<sup>4</sup> canonical reset blocks per seed.

PointGoal2 / trained margin 0.9



| Execution | Return | Hard % | Worst CP95 % | Soft OK |
|-----------|--------|--------|--------------|---------|
| Native    | 17.42  | 4.27   | 7.685        | 1/3     |
| Scale 1   | 16.43  | 1.28   | 4.173        | 3/3     |
| Scale 1.5 | 14.84  | 0.43   | 2.466        | 3/3     |
| Scale 2   | 12.54  | 0.21   | 2.004        | 3/3     |
| Random    | 18.16  | 4.70   | 7.685        | 1/3     |

CarGoal2 / margin schedule to 0.8



| Execution | Return | Hard % | Worst CP95 % | Soft OK |
|-----------|--------|--------|--------------|---------|
| Native    | 19.34  | 2.35   | 5.768        | 1/3     |
| Scale 1   | 16.74  | 0.96   | 2.910        | 1/3     |
| Scale 1.5 | 13.85  | 0.32   | 2.004        | 2/3     |
| Scale 2   | 9.14   | 0.75   | 2.910        | 2/3     |

Filled markers: all three empirical soft budgets met. Worst CP95: largest per-seed nominal one-sided risk bound.  
Exploratory, selected configurations; lines connect measured settings, not a guarantee between them.  
No configuration-search correction. Extra blocks and different evaluator/trainer revisions are kept separate.