

SafeOT execution scale: measured reward / risk trade-offs


Each summary uses training seeds 1, 2, 3 and the same 3 x 10⁴ canonical reset blocks per seed.

PointGoal2 / trained margin 0.9



Execution	Return	Hard %	Worst CP95 %	Soft OK
Native	17.42	4.27	7.685	1/3
Scale 1	16.43	1.28	4.173	3/3
Scale 1.5	14.84	0.43	2.466	3/3
Scale 2	12.54	0.21	2.004	3/3
Random	18.16	4.70	7.685	1/3

CarGoal2 / margin schedule to 0.8



Execution	Return	Hard %	Worst CP95 %	Soft OK
Native	19.34	2.35	5.768	1/3
Scale 1	16.74	0.96	2.910	1/3
Scale 1.5	13.85	0.32	2.004	2/3
Scale 2	9.14	0.75	2.910	2/3

Filled markers: all three empirical soft budgets met. Worst CP95: largest per-seed nominal one-sided risk bound.

Exploratory, selected configurations; lines connect measured settings, not a guarantee between them.

No configuration-search correction. Extra blocks and different evaluator/trainer revisions are kept separate.