



SafeOT execution scale: measured reward / risk trade-offs

Each summary uses training seeds 1, 2, 3 and the same 3×10^4 canonical reset blocks per seed.

PointGoal2 / trained margin 0.9



CarGoal2 / margin schedule to 0.8



Filled markers: all three empirical soft budgets met. Worst CP95: largest per-seed nominal one-sided risk bound.

Exploratory, selected configurations; lines connect measured settings, not a guarantee between them.

No configuration-search correction. Extra blocks and different evaluator/trainer revisions are kept separate.