




Frozen PG2 policies on new resets: 5,616 completed episodes

3 fixed training seeds x 3 execution settings x 624 matched reset IDs per policy



Each line follows one frozen policy. Hard whiskers: exact upper bounds at $\alpha=0.05/3$ per policy. Dashed lines: hard-risk target 2%, empirical soft budget 25. Return/soft points are empirical means. Conditional on these selected policies and reset distribution; no new-training-seed or SOTA claim.